

Computer Memory

Ray Seyfarth

June 29, 2012

Outline

- 1 Memory mapping
- 2 Process memory model in Linux
- 3 Memory example
- 4 Examining memory with gdb
- 5 Examining memory with ebe

Memory mapping

- Computer memory is an array of bytes from 0 to $n - 1$ where n is the memory size
- Programs perceive “logical” addresses which are mapped to physical addresses
- 2 people can run a program starting at logical address 0x4004c8 while using different physical memory
- CPU translates logical addresses to physical during instruction execution
- The CPU translation can be just as fast as if the software used physical addresses
- The x86-64 CPUs can map pages of sizes 4096 bytes and 2 megabytes
- Linux uses 2 MB pages for the kernel and 4 KB pages for programs
- Some recent CPUs support 1 GB pages

Translating an address

- Suppose an instruction references address 0x43215628
- With 4 KB pages, the rightmost 12 bits are an offset into a page
- With 0x43215628 the page offset is 0x628
- The page number is 0x43215
- Let's assume that the computer is set up to translate page 0x43215 to physical addresses 0x7893000 - 0x7893fff
- Then address 0x43215628 is mapped to 0x7893628

Benefits of memory mapping


- User processes are protected from each other
 - ▶ Your vi process can't read my vi's data
 - ▶ Your process can't write my data
- The operating system is protected from malicious or errant code
- It is easy for the operating system to give processes contiguous chunks of “logical” memory

Why study memory mapping?

- If you write programs, the mapping is automatic
- We will not discuss instructions for changing mapping tables
- So what difference does it make?
- It helps explain page faults
 - ▶ Suppose you allocate an array of 256 bytes at logical address 0x45678200
 - ▶ Then all addresses from 0x45678000 to 0x45678fff are valid
 - ▶ You can go well past the end of the array before you can get a segmentation violation
- Knowledge is power!

Process memory model in Linux

- A Linux process has 4 logical segments
 - ▶ text: machine instructions
 - ▶ data: static data initialized when the program starts
 - ▶ heap: data allocated by `malloc` or `new`
 - ▶ stack: run-time stack
 - ★ return addresses
 - ★ some function parameters
 - ★ local variables for functions
 - ★ space for temporaries
- In reality it is more complex
- 131TB is 47 bits of all 1's
- CPU could use 48 bit logical addresses
- Canonical addresses propagate bit 47 through 48-63 so Linux chose to use 47 bits to avoid the top stack address from appearing huge



Memory segments

- The text segment is named `.text` in `yasm`
 - ▶ `_start` and `main` are not actually at 0
 - ▶ The text segment does not need to grow, so the data segment can be placed immediately after it
- The data segment is in 2 parts
 - ▶ `.data` which contains initialized data
 - ▶ `.bss` which contains reserved data (initialized to 0)
 - ▶ “bss” stands for “Block Started by Symbol”
- The heap and the stack both need to grow
 - ▶ The heap grows up
 - ▶ The stack grows down
 - ▶ They meet in the middle and explode
- Use of heap and stack space in assembly does not involve using a named segment

Stack segment limits

- The stack segment is limited by the Linux kernel
- The typical size is 16 MB for 64 bit Linux
- This can be inspected using “`ulimit -a`”
- 16 MB seems fairly small, but it is fine until you start using large arrays as local variables in functions
- The stack address range is 0x7fffffff000000 to 0xffffffffffff
- A fault to addresses in this range are recognized by the kernel to allow the stack to grow as needed

A few adjustments to the memory model

- It appears that the text segment starts at 0x400000 not 0
- Shared libraries map code and data into lots of addresses
- You can map shared memory regions into your programs
- Use “cat /proc/\$\$/maps” to see your shell’s map
 - ▶ \$\$ is the shell’s process id

Memory example source code

```

        segment .data
a        dd      4
b        dd      4.4
c        times   10 dd 0
d        dw      1, 2
e        db      0xfb
f        db      "hello world", 0

        segment .bss
g        resd     1
h        resd     10
i        resb     100
```

Memory example source code (2)

```
segment .text
global  main      ; let the linker know about main
main:
    push    rbp      ; set up a stack frame for main
    mov     rbp, rsp ; set rbp to point to the stack frame
    sub     rsp, 16   ; leave some room for local variables
                    ; leave rsp on a 16 byte boundary
    xor     eax, eax ; set rax to 0 for return value
    leave   ; undo the stack frame manipulations
    ret
```

Memory example listing file

```
1                                %line 1+1 memory.asm
2                                [section .data]
3 00000000 04000000             a dd 4
4 00000004 CDCC8C40             b dd 4.4
5 00000008 00000000<rept>      c times 10 dd 0
6 00000030 01000200             d dw 1, 2
7 00000034 FB                   e db 0xfb
8 00000035 68656C6C6F20776F72-  f db "hello world", 0
9 00000035 6C6400
```

- Addresses are relative to start of .data in this file
- Notice that the 4 byte of 4 is at address 0 (backwards)
- $b = 0x408ccccd = 0\ 10000001\ 00011001100110011001101$
- Sign bit is 0, exponent field is $0x81 = 129$, exponent = 2
- Fraction is $1.00011001100110011001101$

Memory example listing file (2)

```
11                                [section .bss]
12 00000000 <gap>                g resd 1
13 00000004 <gap>                h resd 10
14 0000002C <gap>                i resb 100
```

- Notice that the addresses start again at 0
- The commands reserve space
- `resd 1` reserves 1 double word or 4 bytes
- `resd 10` reserves 10 double words or 40 bytes
- `resb 100` reserves 100 bytes

Memory example listing file (3)

16		[section .text]	
17		[global main]	
18		main:	
19	00000000	55	push rbp
20	00000001	4889E5	mov rbp, rsp
21	00000004	4883EC10	sub rsp, 16
22	00000008	31C0	xor eax, eax
23	0000000A	C9	leave
24	0000000B	C3	ret

Examining memory with gdb

- Time to try some commands in gdb
- Use p for print
 - ▶ Print allows printing expressions
 - ▶ p/d for decimal
 - ▶ try format options t, u, i, c, s, f, a and x
- Examine requires a memory address
 - ▶ x/NFS
 - ▶ N is an optional count
 - ▶ F is a format like print
 - ▶ S is a size character: b=1, h=2, w=4, g=8

Examining memory with ebe

- Run program to a breakpoint
- Control-right-click on a variable name
- Fill in popup form
 - ▶ Variable name
 - ▶ Address - will the `&variable`
 - ▶ Format
 - ★ floating point
 - ★ decimal
 - ★ hexadecimal
 - ★ character
 - ★ string
 - ★ string array (like `argv` in `main`)
 - ▶ Size: 1, 2, 4 or 8 bytes
 - ▶ First and last indices
- Variable will be monitored in data window