

Branching and Looping

Ray Seyfarth

June 29, 2012

Branching and looping

- So far we have only written “straight line” code
- Conditional moves helped spice things up
- In addition conditional moves kept the pipeline full
- But conditional moves are not always faster than branching
- But we need loops to process each bit in a register
- Repeated code can be faster, but there is a limit
- In the next chapter we will work with arrays
- Here we will need to process differing amounts of data
- Repeated code is too inflexible
- We need loops
- To handle code structures like if/else we need both conditional and unconditional branch statements

Outline

- 1 Unconditional jump
- 2 Conditional jump
- 3 Looping with conditional jumps
- 4 Loop instructions
- 5 Repeat string (array) instructions

Unconditional jump

- An unconditional jump is equivalent to a goto
- But jumps are necessary in assembly, while high level languages could exist without goto
- The unconditional jump looks like
`jmp label`
- The `label` can be any label in the program's text segment
- Humans think of parts of the text segment as functions
- The computer will let you jump anywhere
- You can try to jump to a label in the data segment, which hopefully will fail
- The assembler will generate an instruction register (`rip`) relative location to jump
- The simplest form uses an 8 bit immediate: -128 to +127 bytes
- The next version is 32 bits: plus or minus 2 GB
- The short version takes up 2 bytes; the longer version 5 bytes
- The assembler figures this out for you

Unconditional jumps can vary

- An unconditional jump can jump to a location specified by a register's content or a memory location
- You could use a conditional move to hold either of 2 locations in a register and jump to the proper location
- It is simpler to just use a conditional jump
- However you can construct an efficient switch statement by expanding this idea
- You need an array of addresses and an index for the array to select which address to use for the jump

Unconditional jump used as a switch

```
        segment .data
switch: dq      main.case0
        dq      main.case1
        dq      main.case2
i:      dq      2
        segment .text
        global  main                ; tell linker about main
main:   mov     rax, [i]              ; move i to rax
        jmp     [switch+rax*8]       ; switch ( i )

.case0:
        mov     rbx, 100              ; go here if i == 0
        jmp     .end

.case1:
        mov     rbx, 101              ; go here if i == 1
        jmp     .end

.case2:
        mov     rbx, 102              ; go here if i == 2

.end:   xor     eax, eax
        ret
```

Conditional jump

- First you need to execute an instruction which sets some flags
- Then you can use a conditional jump
- The general pattern is
jCC label
- The CC means a condition code

instruction	meaning	aliases	flags
jz	jump if zero	je	ZF=1
jnz	jump if not zero	jne	ZF=0
jg	jump if > zero	jnle ja	ZF=0, SF=0
jge	jump if \geq zero	jnl	SF=0
jl	jump if < zero	jnge js	SF=1
jle	jump if \leq zero	jng	ZF=1 or SF=1
jc	jump if carry	jb jnae	CF=1
jnc	jump if not carry	jae jnb	CF=0

Simple if statement

```
if ( a < b ) {  
    temp = a;  
    a = b;  
    b = temp;  
}
```

```
    mov    rax, [a]  
    mov    rbx, [b]  
    cmp    rax, rbx  
    jge    in_order  
    mov    [temp], rax  
    mov    [a], rbx  
    mov    [b], rax  
in_order:
```


If statement with an else clause

```
if ( a < b ) {  
    max = b;  
} else {  
    max = a;  
}
```

```
        mov    rax, [a]  
        mov    rbx, [b]  
        cmp    rax, rbx  
        jnl    else  
        mov    [max], rbx  
        jmp    endif  
else:    mov    [max], rax  
endif:
```

Looping with conditional jumps

- You could construct any form of loop using conditional jumps
- We will model our code after C's loops
- `while`, `do ... while` and `for`
- We will also consider `break` and `continue`
- `break` and `continue` can be avoided in C, though sometimes the result is less clear
- The same consideration applies for assembly loops as well

Counting 1 bits in a quad-word

```
sum = 0;
i = 0;
while ( i < 64 ) {
    sum += data & 1;
    data = data >> 1;
    i++;
}
```

- There are much faster ways to do this
- But this is easy to understand and convert to assembly

Counting 1 bits in a quad-word in assembly

```
segment .text
global main

main:  mov     rax, [data] ; rax holds the data
      xor     ebx, ebx    ; clear since setc will fill in bl
      xor     ecx, ecx    ; i = 0;
      xor     edx, edx    ; sum = 0;
while: cmp     rcx, 64     ; while ( i < 64 ) {
      jnl     end_while   ; requires testing on opposite
      bt      rax, 0       ; data & 1
      setc    bl          ; move result of test to bl
      add     edx, ebx     ; sum += data & 1;
      shr     rax, 1       ; data = data >> 1;
      inc     rcx         ; i++;
      jmp     while       ; end of the while loop
end_while:
      mov     [sum], rdx   ; save result in memory
      xor     eax, eax     ; return 0 from main
      ret
```

```
    movq    data(%rip), %rax
    movl    $64, %ecx
    xorl    %edx, %edx
```

.L2:

```
    movq    %rax, %rsi
    sarq    %rax
    andl    $1, %esi
    addq    %rsi, %rdx
    subl    $1, %ecx
    jne     .L2
```

- AT&T syntax: operands are reversed and names are more explicit
- The compiler counted down from 64
- Converted the loop to test at the bottom
- Loop has 2 fewer instructions
- Is it faster to use `movq` and `andl`?

Learning from the compiler

- The compiler writers know the instruction set very well
- Most likely `movq` and `andl` is faster
- Testing would tell if the other method is superior
- I also tried the compiler option “`-funroll-all-loops`”
- The compiler added up values for 8 bits in 1 loop iteration
- 8 bits in a 24 instruction loop vs 1 bit in a six instruction loop
- This makes it twice as fast, but the instructions use many different registers allowing parallel execution in 1 core
- Loop unrolling can help a lot with 16 registers
- Examining the generated code should mean than you do no worse
- Clever reorganization can beat the compiler

Do-while loops

- Strict translation of a while loop uses 2 jumps
- It save a jump to the top if you use a do-while loop

```
do {  
    statements;  
} while ( condition );
```

- A do-while loop always executes the loop body at least once
- You can always place an if statement around a do-while to make it behave like a while loop

```
if ( condition ) {  
    do {  
        statements;  
    } while ( condition );  
}
```

- Don't do this in C - let the compiler do it for you

Ugly C code to search through an array

```
i = 0;
c = data[i];
if ( c != 0 ) do {
    if ( c == x ) break;
    i++;
    c = data[i];
} while ( c != 0 );
n = c == 0 ? -1 : i;
```


Assembly code to search through an array

```
mov     bl, [x]           ; value being sought
xor     ecx, ecx          ; i = 0;
mov     al, [data+rcx]    ; c = data[i]
cmp     al, 0             ; if ( c != 0 ) {
jz      end_while        ; skip loop for empty string
```

while:

```
cmp     al, bl            ; if ( c == x ) break;
je      found
inc     rcx               ; i++;
mov     al, [data+rcx]    ; c = data[i];
cmp     al, 0             ; while ( c != 0 );
jnz     while
```

end_while:

```
mov     rcx, -1           ; If we get here, we failed
found:  mov     [n], rcx   ; Assign either -1 or the
                           ; index where x was found
```

Counting loops

```
for ( i = 0; i < n; i++ ) {  
    c[i] = a[i] + b[i];  
}
```

```
        mov     rdx, [n]           ; use rdx for n  
        xor     ecx, ecx           ; i (rdx) = 0  
for:     cmp     rcx, rdx           ; i < n  
        je      end_for           ; get out if equal  
        mov     rax, [a+rcx*8]     ; get a[i]  
        add     rax, [b+rcx*8]     ; a[i] + b[i]  
        mov     [c+rcx*8], rax     ; c[i] = a[i] + b[i];  
        inc     rcx                ; i++  
        jmp     for               ; too bad, loop has 2 jumps  
end_for:
```

- We could use a test before the loop
- We could do loop unrolling

Loop instructions

- The CPU has instructions like `loop` and `loopne` which designed for loops
- They decrement `rcx` and do the branch if `rcx` is not 0
- It is faster to use `dec` and `jnz` instead
- The label must be within -128 to +127 bytes of `rip`
- Probably pointless

```
        mov     ecx, [n]
        sub     ecx, 1
more:    cmp     [data+rcx], al
        loopne  more
        mov     [loc], ecx
```

Repeat string (array) instructions

- The repeat instruction (`rep`) works in conjunction with string (array) instructions
- You first set `rcx` to be the number of repetitions
- You set `rsi` to the address of source data
- And set `rdi` to be the address of destination data
- Then you use a command like

```
rep    movsb
```

- The previous command would copy an array of bytes
- Some string instructions include tests for early termination
- The string instructions can also be used without `rep`

Store instruction

- The `stosb` instruction stores the byte in `al` at the address specified in `rdi` and increments `rdi`
- If the direction flag is set it decrements `rdi`
- There are also `stosw`, `stosd` and `stosq` to operate 2, 4 and 8 byte quantities

```
mov     eax, 1
mov     ecx, 1000000
lea     rdi, [destination]
rep     stosd      ; place 1000000 1's in destination
```

Store instruction

- There are a collection of load string instructions which copy data from the address pointed at by `rsi` and increment (or decrement) `rsi`
- Using `rep lodsb` seems pointless
- The code below uses `lodsb` and optionally `stosb` to copy none carriage return characters

```
        lea    rsi, [source]
        lea    rdi, [destination]
        mov    ecx, 1000000    ; number of iterations
more:    lodsb                    ; get the next byte in al
        cmp    al, 13          ; if al is not 13 store al
        je     skip
        stosb                    ; store al in destination
skip:    sub    ecx, 1          ; count down
        jnz    more
```

Scan instruction

- There are a collection of scan string instructions which scan data from the address pointed at by `rsi` and increment (or decrement) `rsi`
- They compare data against `al`, `ax`, ...
- Below is a version of the C `strlen` function

```
        segment .text
        global strlen

strlen: cld                        ; prepare to increment rdi
        mov     rcx, -1           ; maximum number of iterations
        xor     al, al           ; will scan for 0
        repne   scasb            ; repeatedly scan for 0
        mov     rax, -2          ; start at -1, end 1 past the end
        sub     rax, rcx
        ret
```

Compare instruction

- The compare string instructions compare the data pointed at by rdi and rsi
- The code below implements the C memcmp function

```
segment .text
global  memcmp
memcmp: mov     rcx, rdx
        repe    cmpsb      ; compare until end or difference
        cmp     rcx, 0
        jz      equal      ; reached the end
        movzx   eax, byte [rdi-1]
        movsx   ecx, byte [rsi-1]
        sub     eax, ecx
        ret
equal:   xor     eax, eax
        ret
```


Setting and clearing the direction flag

- The string operations increment their addresses if the direction flag is 0
- They decrement their address if the direction flag is 1
- Use `cld` to prepare for increasing addresses
- Use `std` to prepare for decreasing addresses
- Functions are expected to leave the direction flag set to 0